

## Quiz

---

You are running a company and you want to develop learning algorithms to address each of the following two problems:

Problem 1: You have a large inventory of identical items. You want to predict how many of these items will sell over the next three months.

Problem 2: You would like software to examine individual customer accounts and for each account decide if it has been hacked/compromised.

Should you treat these as classification or as regression problems?

- a) Treat both as classification problems
- b) Treat problem 1 as classification problem, problem 2 as regression problem.
- c) Treat problem 1 as regression problem, problem 2 as classification problem.
- d) Treat both as regression problems

Answer: c)

## Quiz

Of the following examples, which would you address using an unsupervised learning algorithm? (Check all that apply.)

- Given email labeled as spam/not spam, learn a spam filter.
- Given a set of news articles found on the web, group them into set of articles about the same story.
- Given a database of customer data, automatically discover market segments and group customers into different market segments.
- Given a dataset of patients diagnosed as either having diabetes or not, learn to classify new patients as having diabetes or not.

# The Supervised Learning Problem

- Outcome measurement  $Y$  (also called as dependent variable, response, target)
- Vector of  $p$  predictor measurements  $X$  (also called inputs, regressors, covariates, features, independent variables).
- In the regression problem,  $Y$  is quantitative (e.g. price, blood pressure)
- In the classification problem,  $Y$  takes values in a finite, unordered set (survived/died, digit 0-9, cancer class)
- We have training data  $(x_1, y_1), \dots, (x_N, y_N)$ . These are observations (examples, instances) of these measurements.
- On the basis of training data we would like to:
  - Accurately predict unseen test cases
  - Understand which input affect the outcome
  - Asses the quality of our predictions and inferences.

## Unsupervised Learning

- No outcome variable, just a set of predictors (features) measured on a set of examples.
- Objective is more fuzzy – find group of samples that behave similarly, find features that behave similarly, find linear combination of features with the most variation.
- Difficult to know how well you are doing.
- Can be useful as a pre-processing step for supervised learning.


# The key ingredient of machine learning is...

---


## DATA

- Collected from nature or industrial processes.
  - Comes stored in many forms (and formats...):
    - ▶ structured, unstructured, occasionally clean, usually messy, ...
  - In ML we like to view data as a **list of examples** (or we'll turn it into one)
    - ▶ ideally ***many examples of the same nature.***
    - ▶ preferably with each example a ***vector of numbers*** (or we'll first turn it into one!)

# Understanding the ML Task



# Dataset imagined as a point cloud in a high-dimensional vector space



# Machine Learning Tasks (Problem Types)

Supervised learning = predict a target  $y$  from input  $x$

(and semi-supervised learning)

- $y$  represents a category or “class”
  - ➡ classification      binary :  $y \in \{-1, +1\}$  or  $y \in \{0, 1\}$   
multiclass :  $y \in \{1, m\}$  or  $y \in \{0, m - 1\}$
- $y$  is a real-value number
  - ➡ regression       $y \in \mathbb{R}$     or     $y \in \mathbb{R}^m$

Predictive  
models

Unsupervised learning: no explicit prediction target  $y$

- model the probability distribution of  $x$ 
  - ➡ density estimation
- discover underlying structure in data
  - ➡ clustering
  - ➡ dimensionality reduction
  - ➡ (unsupervised) representation learning

Descriptive  
modeling

Reinforcement learning: taking good sequential decisions to maximize a reward in an environment influenced by your decisions.

# Understanding the ML Model

Supervised task:

predict  $y$  from  $x$

input  $x \in \mathbb{R}^d$

target (label)  $y$

n examples

$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$t$
0.32	-0.27	+1	0	0.82	113
-0.12	0.42	-1	1	0.22	34
0.06	0.35	-1	1	-0.37	56
0.91	-0.72	+1	0	-0.63	77
...	...	...	...	...	...

Training set  $D_n$

Learn a function  $f_\theta$  that will minimize prediction errors as measured by cost (loss)  $L$

loss function

$$L(f_\theta(x), y)$$

output  $f_\theta(x)$

$$f_\theta : \text{parameters}$$


$$\begin{matrix} -0.12 & 0.42 & -1 & 1 & 0.22 \end{matrix}$$

$$34$$

input  $x$

target  $y$

# Block Diagram of Learning Machine



## ML Problem Phases

---

The following are the various steps involved in problem solving using ML:

- Define your task
- Collect Data
- Pre-processing of Data
- Dimensionality Reduction/Feature Selection
- Choose ML Algorithm
- Experimental Design
- Test and Validate
- Run System

# Steps used in Machine Learning

- 1. Collecting data:** Be it the raw data from excel, access, text files etc., this step (gathering past data) forms the foundation of the future learning.
  - The better the variety, density and volume of relevant data, better the learning prospects for the machine becomes.
- 2. Preparing the data:** Any analytical process thrives on the quality of the data used.
  - One needs to spend time determining the quality of data and then taking steps for fixing issues such as missing data and treatment of outliers.
- 3. Training a model:** This step involves choosing the appropriate algorithm and representation of data in the form of the model.
  - The cleaned data is split into two parts – train and test (proportion depending on the prerequisites);
  - the first part (training data) is used for developing the model.
  - The second part (test data), is used as a reference.

## Steps used in Machine Learning

---

4. **Evaluating the model:** To test the accuracy, the second part of the data (holdout / test data) is used.
  - This step determines the precision in the choice of the algorithm based on the outcome.
  - A better test to check accuracy of model is to see its performance on data which was not used at all during model build.
5. **Improving the performance:** This step might involve choosing a different model altogether or introducing more variables to augment the efficiency.
  - That's why significant amount of time needs to be spent in data collection and preparation.

# Gradient Descent

- **Gradient Descent** is an optimization algorithm commonly used in machine learning and deep learning to minimize a function, typically a loss or cost function.
- The main idea behind gradient descent is to iteratively adjust the parameters of a model to find the values that minimize the function.
- **How Gradient Descent Works:**
  - **Initialization:** Start with random or pre-defined values for the parameters (weights, biases, etc.) of the model.
  - **Compute the Gradient:** Calculate the gradient of the loss function with respect to the parameters. The gradient is a vector of partial derivatives that points in the direction of the steepest descent.
  - **Update the Parameters:** Move the parameters in the opposite direction of the gradient by a small step (learning rate). This is mathematically expressed as:

$$w^{new} = w^{old} - \alpha \nabla_w J$$

where  $w$ : parameters (weights, biases, etc.)

$\alpha$ : Learning Rate (step size)


$\nabla_w J$ : Gradient of the loss function with respect to the parameters.

- **Repeat:** Continue updating the parameters until the loss function converges (i.e., stops decreasing significantly or reaches a predefined tolerance).

## Gradient Descent (Scalar Case)



- Our task is to improve our guess for  $w$  such that we move from a region of higher gradient to a region of lower gradient.
- For scalar (i.e. one component)  $w$ , this is easy

$$w^{new} = w^{old} - \alpha \frac{dJ}{dw}$$




## Gradient Descent (Vector Case)

- For the vector case we rely on a theorem that says “At any given point the gradient gives the direction of steepest descent”.




## Gradient Descent: Example



$$J(w) = w_1^2 + w_2^2 + 4$$

$$\nabla_w J(w) = \begin{bmatrix} 2w_1 \\ 2w_2 \end{bmatrix}$$



- Gradient Descent gives the iterative formula:

$$w_1^{k+1} = w_1^k - \alpha 2w_1^k$$


$$w_2^{k+1} = w_2^k - \alpha 2w_2^k$$

- We know that the actual minimum is at  $w = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ .
- Let us guess the initial guess is at  $w^0 = \begin{bmatrix} 3 \\ 4 \end{bmatrix}$ .
- Let us see different cases of  $\alpha = 2, 1, 0.1, 0.5$

# Gradient Descent: Example

$$w^0 = \begin{bmatrix} 3 \\ 4 \end{bmatrix}, \alpha = 2$$


Divergent case of  $\alpha$



$$J(w) = w_1^2 + w_2^2 + 4$$

$$\nabla_w J(w) = \begin{bmatrix} 2w_1 \\ 2w_2 \end{bmatrix}$$

$$\begin{aligned} w_1^{k+1} &= w_1^k - \alpha (2w_1^k) \\ w_2^{k+1} &= w_2^k - \alpha (2w_2^k) \end{aligned}$$



Iteration	$w^k$	$\nabla_w J = 2 \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$	$J$	$w^{k+1} = w^k - \alpha \nabla_w J$
0	$\begin{bmatrix} 3 \\ 4 \end{bmatrix}$	$\begin{bmatrix} 6 \\ 8 \end{bmatrix}$	29	$\begin{bmatrix} -9 \\ -12 \end{bmatrix}$
1	$\begin{bmatrix} -9 \\ -12 \end{bmatrix}$	$\begin{bmatrix} -18 \\ -24 \end{bmatrix}$	229	$\begin{bmatrix} 27 \\ 36 \end{bmatrix}$
2	$\begin{bmatrix} 27 \\ 36 \end{bmatrix}$	$\begin{bmatrix} 54 \\ 72 \end{bmatrix}$	2029	$\begin{bmatrix} -84 \\ -108 \end{bmatrix}$

# Gradient Descent: Example

Iteration	$w^k$	$\nabla_w J = 2 \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$	$J$	$w^{k+1} = w^k - \alpha \nabla_w J$	$\alpha = 1$ Oscillates
0	$\begin{bmatrix} 3 \\ 4 \end{bmatrix}$	$\begin{bmatrix} 6 \\ 8 \end{bmatrix}$	29	$\begin{bmatrix} -3 \\ -4 \end{bmatrix}$	
1	$\begin{bmatrix} -3 \\ -4 \end{bmatrix}$	$\begin{bmatrix} -6 \\ -8 \end{bmatrix}$	29	$\begin{bmatrix} 3 \\ 4 \end{bmatrix}$	
2	$\begin{bmatrix} 3 \\ 4 \end{bmatrix}$	$\begin{bmatrix} 6 \\ 8 \end{bmatrix}$	29	$\begin{bmatrix} -3 \\ -4 \end{bmatrix}$	
Iteration	$w^k$	$\nabla_w J = 2 \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$	$J$	$w^{k+1} = w^k - \alpha \nabla_w J$	$\alpha = 0.1$ Converges slowly
0	$\begin{bmatrix} 3 \\ 4 \end{bmatrix}$	$\begin{bmatrix} 6 \\ 8 \end{bmatrix}$	29	$\begin{bmatrix} 2.4 \\ 3.2 \end{bmatrix}$	
1	$\begin{bmatrix} 2.4 \\ 3.2 \end{bmatrix}$	$\begin{bmatrix} 4.8 \\ 6.4 \end{bmatrix}$	20	$\begin{bmatrix} 1.92 \\ 2.56 \end{bmatrix}$	
2	$\begin{bmatrix} 1.92 \\ 2.56 \end{bmatrix}$	$\begin{bmatrix} 3.84 \\ 5.12 \end{bmatrix}$	14.24	$\begin{bmatrix} 1.536 \\ 2.048 \end{bmatrix}$	
30	$\begin{bmatrix} 0.0037 \\ 0.005 \end{bmatrix}$	...	4.000	...	
Iteration	$w^k$	$\nabla_w J = 2 \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$	$J$	$w^{k+1} = w^k - \alpha \nabla_w J$	$\alpha = 0.5$ Converges rapidly
0	$\begin{bmatrix} 3 \\ 4 \end{bmatrix}$	$\begin{bmatrix} 6 \\ 8 \end{bmatrix}$	29	$\begin{bmatrix} 0 \\ 0 \end{bmatrix}$	
1	$\begin{bmatrix} 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \end{bmatrix}$	4	$\begin{bmatrix} 0 \\ 0 \end{bmatrix}$	

## Some Lessons from the example

---

- It is possible for gradient descent algorithm to:
  - Diverge ( $\alpha = 2$ )
  - Oscillates without diverging or converging ( $\alpha = 1$ )
  - Converge slowly ( $\alpha = 0.1$ )
  - Converge rapidly ( $\alpha = 0.5$ )
- The choice of  $\alpha$  is important and is part of algorithm design.
- $\alpha$  is a **hyper parameter** (a parameter that must be set before learning begins).

## Stopping Criteria of Gradient Descent


- Ideally, we should stop when  $\nabla_w J = 0$ .
- This almost never happens in practice as the number of iterations required could be infinite.
- In practice, we decide on some precision  $\epsilon$  (say  $\epsilon \approx 10^{-5}$ )
- Multiple options for stopping criteria. Stop when
  - $\|w^{k+1} - w^k\| \leq \epsilon$
  - $\|\nabla_w J(w^k)\| \leq \epsilon$
  - $\|J(w^{k+1}) - J(w^k)\| \leq \epsilon$

## Summary of Gradient Descent Procedure

---

1. Decide on  $\alpha, \epsilon$  and stopping criterion.
2. Make an initial guess for  $w = w^0$ .
3. Calculate  $w^{k+1} = w^k - \alpha \nabla_w J$ 
  - a) Calculate gradient numerically, if required.
4. Calculate stopping criterion.
  - a) If satisfied, stop
  - b) If not satisfied, go to Step 3.

# The Learning Paradigm




- We wish to learn the relationship between the input and output data.
- For now, we will think of this relationship as a function (which relates  $x \rightarrow y$ ).
  - We call this function the **model** or **hypothesis function**.
- The function has two parts:
  - Form of the function [e.g.  $h(x) = w_0 + w_1x_1 + w_2x_2 + w_3x_1x_2$ ]
  - Parameter of the function
- Typically machine learning **learns only the parameters**.
- The form is provided by **ML Engineer**, which requires **domain knowledge**.

## Forward Modeling



- A model or hypothesis is an educated guess at what the relationship between input and output is.
- As mentioned before, it has two pieces:
  - Form of the function: Linear, Quadratic, Exponential etc.
  - Parameters of the function
- We sometimes use the notation  $y = f(x; w)$ 
  - Given  $x$  and choice of  $w$ , we can find a corresponding  $y$ .
- The function  $f$  going from  $x$  to  $y$  is called the forward model.
  - The process is sometimes called feedforward

# Learning the parameters via feedback



To learn the parameters we follow this paradigm.

- Collect lots of data pairs (Input Vector, Output Vector) =  $(x, y)$
- Guess for the form of the hypothesis function  $h(x; w)$ 
  - Example:  $h(x; w) = w_0 + w_1 x_1 + w_2 x_2$
- For an arbitrary guess for  $w$ 
  - We will get some  $\hat{y} = h(x; w)$  which will not match the ground truth  $y$ .
- Define a cost function  $J(y, \hat{y}(w))$  depending on the difference.
- Find the optimal  $w$  by minimizing  $J(w)$ .
  - By using some optimization procedure such as gradient descent

## What the ML Engineer must provide

- Appropriate decisions for input and output vectors  $(x, y)$ 
  - All problems are data, all solutions are functions/maps
- Choosing appropriate datasets (should be usually large).
- Some appropriate form of the forward model  $\hat{y} = h(x; w)$ 
  - Example: Linear Model  $\hat{y} = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$
- Form the loss function
  - Example: Least Squares  $J(y, \hat{y}) = (y - \hat{y})^2$
- Optimization Algorithm – Example: Gradient Descent
  - And associated hyper parameters such as  $\alpha$ .

Machine Learning is not magic! It requires a lot of input from engineers.

## A look ahead in this course

---


- We will be looking in the next few weeks at various types of hypothesis function  $\hat{y} = h(x; w)$ .
- Each of these have their own purpose and domain where they work well:
  - **Linear Regression:** For simple polynomial regression problems
  - **Logistic Regression:** For two-class (binary) classification problems
  - **Deep Neural Networks:** For any general non-linear problems
    - ▶ There is a theorem that assures us that sufficiently large neural network will approximate any function.
- There are also possible loss functions for each. We will see what is the loss function for each hypothesis function.
- It is possible you might have a better model than these for your problems. However, the general procedure outlined here remains the same.

# Regression

- Regression is a statistical method used for predicting continuous outcome (dependent variable) based on one or more input features (independent variables).
- Purpose of regression is to identify the relationship between variables, understand patterns and make future predictions.
- **Goal:** Given a training set comprising of  $N$  observations  $\{x_i\}, i = 1, 2, \dots, N$ , together with corresponding target values, the goal is to predict the value of  $y$  for a new value of  $x$ .

# Types of Regression Models


1. **Simple Linear Regression:** Relationship between two variables.
2. **Multiple Linear Regression:** Relationship between one dependent and multiple dependent variables.
3. **Polynomial Regression:** Extension of Linear Regression by adding polynomial terms.



# Introduction to Linear Regression


- Problem: Given  $m$  number of data points  $(x_i, y_i)$ , where  $x_i$  is the **size of a house** in square feet and  $y_i$  is its **price** in 1000 \$s. We are asked to predict the price of a house (not in the data set), given its size  $x$ .

Housing price prediction.



# Linear Regression


- Regression is a method of modelling a target value based on independent predictors.
- Simple linear regression is a type of regression analysis where the number of independent variables is one and there is a linear relationship between the independent ( $x$ ) and dependent ( $y$ ) variable.
- Based on the given data points, we try to plot a line that models the points the best.
- The line can be modelled based on the linear equation:  $y = mx + c$
- The motivation of the linear regression algorithm is to find the best values for  $m$  and  $c$ .



# Linear Regression

- Let us write the equation  $y = mx + c$  as  

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$
- In short  $h: x \rightarrow y$ , where  $h$  is known as the hypotheses and  $\theta_0, \theta_1$  are the parameters.
- So, we are interested to find the best values of the parameters  $\theta_0, \theta_1$  so that the hypothesis fits the data points the best.




## Cost Function

- The cost function helps us to figure out the best possible values for  $\theta_0$  and  $\theta_1$ , which would provide the best fit line for the data points.
- Choose  $\theta_0$ ,  $\theta_1$  such that  $h_\theta(x)$  is close to  $y$  for our training examples  $(x, y)$ .
- Minimize  $\frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$   
where  $m$ : number of training examples and  $h_\theta(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$
- Formally the cost function (least squared error function) is:  
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

# Cost Function

- We choose the above function to minimize.
- The difference between the predicted values and ground truth measures the error difference.
- We square the error difference and sum over all data points and divide that value by the total number of data points.
- This cost function is also known as the **Mean Squared Error (MSE)** function as it provides the average squared error over all the data points.
- Now, using this MSE function we are going to change the values of  $\theta_0$  and  $\theta_1$  such that the MSE value settles at the minima.



# Cost Function: Intuition I

- In formal way


- Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x$
- Parameters:  $\theta_0, \theta_1$
- Cost Function:  $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$   
where  $h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$
- Goal:  $\underset{\theta_0, \theta_1}{\text{Minimize}} J(\theta_0, \theta_1)$

- A simplified Hypothesis Function

- Hypothesis:  $h_{\theta}(x) = \theta_1 x$  [ $\theta_0 = 0$ ]
- Parameters:  $\theta_1$
- Cost Function:  $J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  where  $h_{\theta}(x^{(i)}) = \theta_1 x^{(i)}$
- Goal:  $\underset{\theta_1}{\text{Minimize}} J(\theta_1)$

## Cost Function: Intuition I

- Comparison of Hypothesis ( $h_{\theta}(x)$ ) with Cost Function ( $J(\theta_1)$ ):
  - The hypothesis is a function of  $x$  for fixed  $\theta_1$  and cost function is a function of the parameter  $\theta_1$
- Let us assume the training set contains training data  $(1,1)$ ,  $(2,2)$  and  $(3,3)$ . So  $m = 3$ .






# Cost Function: Intuition I

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 = \frac{1}{2m} \sum_{i=1}^m (\theta_1 x^{(i)} - y^{(i)})^2$$

$$J(1) = \frac{1}{2 \times 3} [(1-1)^2 + (2-2)^2 + (3-3)^2] = 0$$


$$J(0.5) = \frac{1}{2 \times 3} [(0.5-1)^2 + (1-2)^2 + (1.5-3)^2] = 0.58$$

$$J(0) = \frac{1}{2 \times 3} [(0-1)^2 + (0-2)^2 + (0-3)^2] = 2.3$$





## Cost Function: Intuition I

- We have computed  $J(1) = 0$ ,  
 $J(0.5) = 0.58$ ,  $J(0) = 2.3$ .
- In the similar way we can compute  
 $J(-0.5) = 5.25$ ,  $J(1.5) = 0.58$ ,  $J(2) = 2.3$   
and  $J(2.5) = 5.25$ .
- Plotting  $(\theta_1, J(\theta_1))$  we have the following graph:
- Optimal value of  $\theta_1$  is 1.



## Cost Function: Intuition II

**Contour Plots:** Each of the ovals shows the set of points takes the same value of  $J(\theta_0, \theta_1)$



## Least Square Method of Solving Simple Linear Regression

- For  $i^{\text{th}}$  data point the predicted value  $\hat{y}^{(i)} = \theta_0 + \theta_1 x^{(i)}$ .
- The error (residual) for  $i^{\text{th}}$  data point:  $Error_i = y^{(i)} - \hat{y}^{(i)} = y^{(i)} - (\theta_0 + \theta_1 x^{(i)})$ .
- The sum of squared error of the model,

$$SSE(\theta_0, \theta_1) = \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2 = \sum_{i=1}^N (y^{(i)} - (\theta_0 + \theta_1 x^{(i)}))^2$$

- Partial derivative of SSE w.r.t,  $\theta_1$ :

$$\frac{\partial(SSE)}{\partial\theta_1} = \frac{\partial}{\partial\theta_1} \sum_{i=1}^N (y^{(i)} - (\theta_0 + \theta_1 x^{(i)}))^2 = -2 \sum_{i=1}^N x^{(i)}(y^{(i)} - \theta_0 - \theta_1 x^{(i)})$$

$$\frac{\partial(SSE)}{\partial\theta_1} = 0 \Rightarrow -2 \sum_{i=1}^N x^{(i)}(y^{(i)} - \theta_0 - \theta_1 x^{(i)}) = 0 \Rightarrow \sum_{i=1}^N x^{(i)}y^{(i)} = \theta_1 \sum_{i=1}^N (x^{(i)})^2 + \theta_0 \sum_{i=1}^N x^{(i)}$$

----- (1)

## Least Square Method of Solving Simple Linear Regression

- Partial derivative of SSE w.r.t,  $\theta_0$ :

$$\frac{\partial(SSE)}{\partial\theta_0} = \sum_{i=1}^N (-1)2(y^{(i)} - \theta_0 - \theta_1x^{(i)}) = -2 \sum_{i=1}^N (y^{(i)} - \theta_0 - \theta_1x^{(i)})$$

$$\frac{\partial(SSE)}{\partial\theta_0} = 0 \Rightarrow -2 \sum_{i=1}^N (y^{(i)} - \theta_0 - \theta_1x^{(i)}) = 0 \Rightarrow \sum_{i=1}^N y^{(i)} = \theta_1 \sum_{i=1}^N x^{(i)} + \theta_0 \sum_{i=1}^N 1$$

$$\Rightarrow \sum_{i=1}^N y^{(i)} = \theta_1 \sum_{i=1}^N x^{(i)} + \theta_0 N \quad \text{----- (2)}$$

- From equation (2) we have:  $\theta_0 = \frac{\sum y^{(i)} - \theta_1 \sum x^{(i)}}{N}$  ----- (3)

- Substituting the value of  $\theta_0$  in equation (1):  $\sum x^{(i)}y^{(i)} = \theta_1 \sum (x^{(i)})^2 + \left(\frac{\sum y^{(i)} - \theta_1 \sum x^{(i)}}{N}\right) \sum x^{(i)}$
- $$\Rightarrow \theta_1 = \frac{\sum x^{(i)} \sum y^{(i)} - N \sum x^{(i)}y^{(i)}}{(\sum x^{(i)})^2 - N \sum (x^{(i)})^2} \Rightarrow \theta_1 = \frac{N \sum x^{(i)}y^{(i)} - \sum x^{(i)} \sum y^{(i)}}{N \sum (x^{(i)})^2 - (\sum x^{(i)})^2}$$

## Least Square Method of Solving Simple Linear Regression

### Example:

Consider the following dataset of points:

x	y
1	2
2	3
3	5
4	7
5	8

### Solution:

$$\theta_1 = \frac{5 \times 91 - 15 \times 25}{5 \times 55 - (15)^2} = 1.6$$

$$\theta_0 = \frac{25 - 1.6 \times 15}{5} = 0.2$$

Find the best fit line using the list square method.

$$\theta_1 = \frac{N \sum x^{(i)}y^{(i)} - \sum x^{(i)} \sum y^{(i)}}{N \sum (x^{(i)})^2 - (\sum x^{(i)})^2}$$

$$\theta_0 = \frac{\sum y^{(i)} - \theta_1 \sum x^{(i)}}{N}$$